Decision Log

Arrival — what was cut, and why

Kevin Francisco · September 2026

A PRD says what the product is. This says what it is not, and what each choice cost. Decisions are listed in the order they were made.

1. Scope to one trip, not to fluency

Chose. The product ends when the user lands. It teaches a finite set of items sized to a specific trip and never claims to teach a language.

Why. Existing apps are built for open-ended fluency and have no concept of a departure date. A traveler on the Fodor's forums described the available programs as too general for their needs, which is the gap stated in a user's own words.

Gave up. Retention. A fluency product can hold a user for years; this one is designed to be finished and set down. Also rules out the language-hobbyist segment entirely.

2. The scheduler is the product, not the content

Chose. Investment goes into the plan generator: sizing the plan to available time and clamping spaced repetition so peak recall lands on the trip dates.

Why. Audio, offline access, roleplay, and culture notes are table stakes that TripLingo shipped a decade ago. Deadline-aware scheduling is the one mechanic no competitor has, and the one thing a general chatbot cannot hold across sessions.

Gave up. Content depth at launch. The item library is shallower than any incumbent's.

3. Rank by traveler regret, not by frequency

Chose. Items are ranked partly by how many travelers to that city said they wished they had known them, and every item shows why it ranked where it did.

Why. The strongest objection to a phrasebook is that you do not know what matters until you are in the situation. The answer is to let people who were already there do the ranking. It is also the only proprietary data asset available in this category.

Gave up. Cold-start quality. The corpus must be seeded manually before it carries any real signal.

4. No streaks, XP, leaderboards, or hearts

Chose. The departure countdown is the only motivator. No gamification layer.

Why. Gamification is a retention mechanic built for open-ended goals. Duolingo cut monthly churn from ~47% to ~28% with it, but it is solving a problem this product does not have. A hard date already supplies the pressure.

Gave up. A proven engagement lever, and the daily-habit hook that drives most category retention benchmarks.

5. Comprehension drills over more vocabulary

Chose. Each question item also trains the three or four replies the traveler is likely to receive.

Why. The most under-served pain in the category. Every phrasebook trains production; none train receiving the answer at native speed. Being unable to parse the reply is what actually makes travelers freeze.

Gave up. Roughly half the item budget in each plan, which could otherwise have gone to more phrases.

6. Intake before signup

Chose. Five intake questions, plan reveal, and the first lesson all happen anonymously. The account prompt comes after lesson one.

Why. The plan reveal is the only moment that differentiates this from the six apps a user has already abandoned. Asking for an email before showing it means asking before earning it. Intake is only 90 seconds, so the cost of anonymity is low.

Gave up. Anonymous bounces are unrecoverable and cannot be re-engaged by email. Requires local persistence so a refresh does not destroy the plan.

7. Plan re-fit instead of failure

Chose. When a user falls behind pace, the app shrinks and re-ranks the plan rather than letting them fail against it.

Why. Falling behind and then abandoning is the dominant failure mode in this category. Re-fit also makes the scheduling engine visible to the user a second time, which is the thing worth showing off.

Gave up. Plan integrity. A re-fit plan is a weaker plan, and repeated re-fits could let a user coast to a nearly empty deck.

8. Three languages, nine cities, city-level content

Chose. Japan (Tokyo, Kyoto, Osaka), Germany (Berlin, Munich, Hamburg), Costa Rica (San José, La Fortuna, Tamarindo). Content is written per city, and Costa Rican Spanish uses local usage rather than textbook Spanish.

Why. Depth demonstrates the thesis better than breadth. The useful phrase set genuinely differs between Osaka street stalls and Kyoto temple entry, and showing that is the proof that trip-scoping is real rather than cosmetic.

Gave up. Most of the addressable market at launch. Anyone traveling anywhere else cannot use the product.

9. Mascot tied to readiness, not to streak

Chose. A caterpillar-to-butterfly arc driven by trip readiness percentage. It completes at 100% and stops. It never appears when the user is behind.

Why. Metamorphosis is a finite arc with an endpoint, which reinforces the “this ends when you land” framing instead of fighting it. A mascot that shows up while you are failing becomes a nag, and the calm tone collapses.

Gave up. The daily-return pressure a nagging mascot creates.

10. Web only, TTS audio, no native app

Chose. Mobile-responsive web, synthesized audio, no app store presence at launch.

Why. The MVP validates one hypothesis: that a plan sized to a real deadline gets finished when generic courses do not. Native speaker recordings and app store distribution do not affect that test.

Gave up. Audio quality, offline reliability, push notifications, and app store discovery.

11. Measure plan completion by departure date, nothing else

Chose. Primary metric is the share of users who finish their plan before they leave. DAU, session length, and total items learned are explicitly ignored.

Why. Those are engagement metrics for an open-ended product. This one succeeds by being finished and set down, so optimizing for time-in-app would push the product toward the thing it is defined against.

Gave up. Comparability with category benchmarks, and the metrics most investors and reviewers expect to see.

12. Positioned as a portfolio piece and a wedge, not a venture

Chose. Built to demonstrate the scheduling engine. Not positioned as a standalone consumer subscription business.

Why. The economics of a one-trip consumer subscription are structurally weak: narrow acquisition window, one-purchase lifetime value, a free general-purpose substitute, and a category monopoly. TripLingo ran this experiment for eight years and exited into corporate travel.

Gave up. Nothing yet. This is a positioning decision, revisited if the completion metric comes back strong.

Open questions

- What the learn-ahead ceiling should be, if any, once a user pulls items forward repeatedly.
- Whether the post-trip regret prompt should be email, in-app, or both.
- Whether per-trip pricing or an embedded B2B placement is the first monetization test.