

# Product Requirements Document

**Product:** Arrival (working name) **Category:** Trip-scoped language preparation **Version:** 1.0 (MVP definition)  
**Date:** September 5, 2026 **Author:** Kevin Francisco

---

## Product Overview

Arrival is a trip-scoped language preparation product. A traveler enters their destination, dates, planned activities, and available study time. Arrival returns a finite, dated learning plan that ends when their flight leaves, covering only the language they are likely to need on that specific trip.

The product is explicitly not a language course. It does not teach a language. It prepares a traveler for a defined set of real conversations that will occur on defined days, and it tells them how ready they are.

**Problem being solved.** Existing products are open-ended curricula with no finish line and no awareness of a departure date. Travelers report studying for months and being unable to speak on arrival. Content is not ranked by real-world travel frequency: the widely cited example is Duolingo's opening Romanian lesson teaching "I am a boy" and "I am a man," phrases no traveler needs. Meanwhile the highest-value need goes unmet entirely. As one traveler put it on a Fodor's forum thread about learning Spanish for a month in Malaga and Seville, the problem with knowing essential phrases is that once you ask the question, the person answers in Spanish and you do not understand the reply.

### Three product commitments that follow from the research:

1. **Finishable.** Every plan has an end date and a stated ceiling. The user is told up front what they will and will not be able to do by departure.
  2. **Bidirectional.** Every scenario trains what the user says *and* the replies they will hear back, with destination-region audio at native speed.
  3. **Scoped by trip, priced by trip.** A one-time trip pass, not a subscription. Education apps see over 95% of users churn by day 30 and roughly 2% day-30 retention, among the lowest of any app category. A subscription attached to a 3-to-8-week product lifecycle produces cancellations and negative reviews by design.
- 

## Target Audience

### Primary user: the Committed Short-Trip Traveler

| Attribute      | Profile                                                                |
|----------------|------------------------------------------------------------------------|
| Age            | 28 to 55                                                               |
| Trip type      | 7 to 21 day leisure international trip, booked 3 weeks to 6 months out |
| Language level | True beginner or lapsed beginner in the destination language           |

---

| Attribute      | Profile                                                                    |
|----------------|----------------------------------------------------------------------------|
| Time available | 10 to 20 minutes per day, inconsistently                                   |
| Spend behavior | Already buys eSIMs, travel insurance, lounge passes, skip-the-line tickets |
| Trigger moment | Flights booked and confirmed                                               |

**What they want:** to not feel helpless, to not be the tourist who cannot say anything, to handle a restaurant, a train, and a hotel without panic-translating.

**What they do not want:** fluency, grammar instruction, a streak, or a curriculum that outlives their trip.

### Secondary user: the Anxious First-Timer to a High-Barrier Destination

Traveling to a country with low English fallback and a non-Latin script (Japan, Korea, Thailand, Vietnam). Higher anxiety, higher willingness to pay, weaker substitutes because camera translation of a menu is slower and less pleasant than knowing the six words that matter.

### Explicitly out of scope for v1

Learners pursuing fluency, students, expats and long-stay relocators, business travelers, and anyone without a booked trip. Each of these breaks the finish-line premise the product is built on.

### Needs the research surfaced, mapped to who feels them

| Need                                            | Evidence                                                    | Severity |
|-------------------------------------------------|-------------------------------------------------------------|----------|
| Understand the reply, not just ask the question | Recurring forum complaint; no competitor solves it          | Critical |
| Know what is actually worth learning            | Forum poster: most programs are "too general for our needs" | Critical |
| Know if the time available is enough            | No competitor calibrates to a departure date                | High     |
| Regionally correct vocabulary                   | Traveler drilled "abacaxi," Portugal uses "ananas"          | Medium   |
| Retain it until the trip, not lose it before    | Decay between study and departure                           | High     |
| Recover mid-conversation when lost              | Rescue phrases are buried in every product                  | High     |

## Core Features

### Trip Plan Generator

Converts trip parameters into a finite, ranked, dated curriculum.

**Inputs:** destination city, departure date, return date, activity selections (restaurants, hotels, transit, sightseeing, shopping, nightlife, hiking, emergencies), self-reported level, minutes available per day.

**Logic:** score every scenario in the library on expected encounters per trip x cost of failing x learnability. Fill the plan to the user's realistic capacity, calculated as days until departure x minutes per day x retention factor. Cut everything below the line and show the user what was cut and why.

**Output:** an ordered list of scenario modules with a per-day schedule and an explicit capacity statement, for example: "21 days at 15 min/day gets you 9 scenarios. You will handle restaurants, trains, and shops. You will not handle open-ended small talk."

**Why this feature:** the two most cited frustrations are content that is too general and no sense of what is achievable in the time available. This is the only feature that addresses both, and it is the product's core differentiator against Babbel's fixed CEFR path.

### **Bidirectional Scenario Modules**

Each module trains one real interaction in both directions.

#### **Structure per module:**

- 4 to 8 phrases the user will say
- 5 to 8 replies the user will hear, recorded in the destination region at native speed
- A listening drill where the user identifies the reply, not just repeats a line
- One rescue line drilled in every single module ("I did not understand, slowly please")
- A short cultural note where getting it wrong is socially costly

**Why this feature:** this is the defensible content asset. Memrise comes closest with native-speaker video clips that prepare your ears for real speed and accents, but does not tie them to a scenario script. Nothing else trains comprehension of the specific replies a specific scenario produces. An LLM can generate the phrase list on demand for free; it cannot generate a regionally accurate recorded audio corpus.

### **Departure-Anchored Readiness Engine**

Spaced repetition scheduled backward from the departure date, plus a visible readiness score.

#### **Mechanics:**

- Review intervals are computed so peak recall lands on arrival day, not on lesson completion
- A taper in the final 5 days: fewer new items, more retrieval of weak items
- A departure-eve rehearsal that runs the full trip in order
- Readiness score per scenario (0 to 100) and one composite trip score, visible on the home screen every session

**Why this feature:** open-ended curricula have no target date, so decay between study and travel is unmanaged. The readiness score is also the product's core motivational mechanic, replacing streaks. A streak measures attendance. A readiness score measures the thing the user actually bought.

# Essential User Journeys

## Onboarding: booking to first plan

**Target: under 90 seconds from open to first lesson. No paywall before the plan is visible.**

1. **Destination.** City-level, not country-level. Enables regional vocabulary and audio selection.
2. **Dates.** Departure and return. Computes both trip length and days remaining.
3. **Activities.** Multi-select from 8 icons. No free text in v1.
4. **Level.** Three options: none at all, a few words, I studied it once.
5. **Daily time.** Three options: 5, 15, or 30 minutes.
6. **Plan reveal.** Full plan shown at once, with the capacity statement and the explicit cut list.
7. **First lesson starts immediately.** Paywall appears after the first scenario completes, once value has been demonstrated.

**Design rationale:** onboarding length and time-to-value are the dominant day-1 retention drivers, and a paywall placed before the personalized plan is the single most common failure pattern in edtech onboarding. The plan itself is the hook, so it must be free to see.

## Learning: the daily session

Sessions are fixed to the user's declared minute budget and never run over.

Home screen

- > "Today: Ordering at a restaurant, 15 min"
- > Say-side drill (phrases the user produces)
- > Hear-side drill (replies the user must recognize)
- > Rescue line, every session
- > Mixed retrieval of 3 to 5 weak items from prior modules
- > Readiness score updates, visible immediately
- > Session ends. No "one more lesson" prompt.

### Rules:

- Never surface content outside the plan
- Never extend a session past the stated time budget
- No streak, no hearts, no lives, no loss framing
- If a user misses days, the plan silently recalculates and re-cuts to the new capacity rather than accumulating a backlog

**Design rationale:** overwhelm and guilt are recurring reasons people abandon language apps. A plan that shrinks honestly when the user falls behind is more useful than one that accrues debt.

## Progress and completion: readiness, not completion percentage

**Home screen shows three things and nothing else:**

- Days until departure

- Composite readiness score
- Today's session

**Readiness display:** a per-scenario grid, each scenario colored by readiness band (Not started / Shaky / Ready). The user should be able to see at a glance that they are ready for restaurants and shaky on trains.

#### Milestones:

- **T-7 days:** plan locks. No new scenarios added, retrieval only.
- **T-1 day:** departure-eve rehearsal, full trip run-through in itinerary order.
- **Departure day:** app switches to Trip Mode.

**Trip Mode:** the drilled content only, offline, ordered by scenario, with the rescue line pinned to the top of every screen. No lessons, no scoring, no new material. This is the phase where competitors hand the user off to Google Translate, and where the product should feel like the payoff for the prior three weeks.

**Post-trip (day +3):** a single prompt asking which scenarios they actually used and what came up that the plan missed. This is the primary content improvement loop and the primary source of referral.

## Design Guidelines

### Principles

1. **Countdown, not streak.** The organizing visual element is time until departure. Every screen should make the deadline felt. This replaces gamification entirely.
2. **Finite by design.** The user must always be able to see the whole plan, start to finish, on one screen. No infinite scroll, no locked future path, no hidden content.
3. **Honest capacity.** Never imply a level of readiness the user does not have. Under-promising is a feature; the product's credibility rests on the readiness score being trustworthy.
4. **Travel utility, not education.** The product should feel like an eSIM or a boarding pass, not a classroom. Nearest reference points: airline apps and travel wallets, not Duolingo.
5. **Audio-first.** Every phrase and every reply is playable in one tap, with a slow-speed variant. The listening side must feel as prominent as the speaking side, because it is the differentiator.
6. **Offline as a first-class state.** Trip Mode must work with zero connectivity. Design and test the offline state before the online one.

### Visual direction

| Element | Direction                                                                                                                                  |
|---------|--------------------------------------------------------------------------------------------------------------------------------------------|
| Tone    | Calm, practical, adult. No mascot, no cartoon characters, no anthropomorphized coach                                                       |
| Palette | Neutral base with one accent color for readiness state. Readiness bands use a three-step scale, avoiding red-green only, for accessibility |

| Element    | Direction                                                                                                                              |
|------------|----------------------------------------------------------------------------------------------------------------------------------------|
| Typography | High-legibility sans for UI. Destination script rendered large, with romanization always available and always toggleable, never forced |
| Density    | One primary action per screen. Sessions are one card at a time                                                                         |
| Motion     | Minimal. Reserve animation for the readiness score updating, which is the single moment of reward in the product                       |
| Imagery    | Destination photography on the plan screen only. Everywhere else, text and audio                                                       |

### Anti-patterns to avoid

- Loss aversion mechanics (streaks, hearts, "your streak is at risk")
- Push notifications framed as guilt; frame as countdown ("14 days out, 15 minutes today")
- Any content recommendation outside the trip plan
- Celebrating lesson completion rather than readiness gain

## MVP Scope

### In scope for v1

| Item                   | Detail                                                                              | Rationale                                                                                                                         |
|------------------------|-------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------|
| One language           | Japanese, English speakers only                                                     | Low English fallback, script barrier weakens translation-app substitutes, high traveler anxiety, strong US and EU outbound volume |
| One destination tier   | Tokyo, Kyoto, Osaka                                                                 | City-level regional accuracy is a differentiator; three cities keeps the audio corpus tractable                                   |
| 8 activity categories  | Restaurants, hotels, transit, sightseeing, shopping, nightlife, hiking, emergencies | Matches the activity taxonomy travelers already think in                                                                          |
| ~24 scenario modules   | Bidirectional, recorded audio                                                       | Enough to fill a 6-week plan at 30 min/day, or a 2-week plan at 15 min/day                                                        |
| Trip Plan Generator    | Rule-based scoring, not ML                                                          | Deterministic and explainable at this scale; ML is unjustified until there is usage data                                          |
| Departure-anchored SRS | Backward-scheduled, with T-7 lock and T-1 rehearsal                                 | Core mechanic                                                                                                                     |
| Readiness scoring      | Per scenario and composite                                                          | Core mechanic and core motivation                                                                                                 |
| Trip Mode              | Offline, itinerary-ordered, rescue line pinned                                      | The in-destination payoff                                                                                                         |

| Item             | Detail                                                       | Rationale                                                                              |
|------------------|--------------------------------------------------------------|----------------------------------------------------------------------------------------|
| Post-trip survey | Single screen, day +3                                        | Content loop and referral loop                                                         |
| Pricing          | \$24 one-time trip pass, access through return date + 7 days | Matches trip lifecycle and traveler spend behavior; avoids the subscription-churn trap |

### Explicitly cut from v1

| Cut                           | Reason                                                                                                                                           |
|-------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------|
| Open-ended AI voice roleplay  | Speak raised \$162M at a \$1B valuation and does this better; Duolingo Max ships Video Call and Roleplay. Competing here is expensive and losing |
| Streaks, XP, gamification     | Contradicts the finite-plan premise and the target user's stated preferences                                                                     |
| Grammar instruction           | Travelers consistently advised to skip grammar drills and focus on set phrases                                                                   |
| Additional languages          | Each language is a full audio corpus. Do not scale content before validating willingness to pay                                                  |
| Reading and writing           | Beyond menu and sign recognition, irrelevant to a 7-day trip                                                                                     |
| Itinerary file import         | High effort, unproven value. Manual activity selection tests the same hypothesis                                                                 |
| Social, leaderboards, sharing | No mechanism to drive them in a single-trip product                                                                                              |
| Web app                       | Mobile-first; the product is used on a phone in a restaurant                                                                                     |

### Future roadmap, in priority order

1. **Booking-moment distribution.** Partnerships with OTAs, airlines, and tour operators to trigger the offer at the moment flights are confirmed. This is the highest-leverage item on the roadmap and arguably the actual business. TripLingo built a comparable consumer product, raised \$2.3M, reached 500,000+ travelers, and ultimately found its value inside a corporate travel management platform after being acquired in 2019. Start where they finished.
2. **Language expansion.** Korean, Thai, Vietnamese, Italian, in that order. Prioritize low-English-fallback destinations over high-volume ones.
3. **Itinerary ingestion.** Parse booking confirmations to auto-select activities and order Trip Mode by day.
4. **Scenario-scoped voice practice.** Only within a drilled scenario, never open-ended. Adds speaking confidence without competing head-on with Speak.
5. **Companion plans.** Shared plan for two travelers with split scenario coverage, one takes transit, one takes restaurants.
6. **Repeat-trip pricing.** A multi-trip pass once repeat purchase behavior is observed, not before.

## Validation Gates Before Building

Do not build v1 until these are answered. The cheapest test is a landing page plus a hand-built plan delivered to 30 travelers who have already booked Japan flights.

| Gate                                | Metric                                                      | Threshold |
|-------------------------------------|-------------------------------------------------------------|-----------|
| Willingness to pay at the trigger   | Paid conversion from booked-trip traffic                    | > 8%      |
| The finish-line premise holds       | Share completing 70%+ of plan before departure              | > 40%     |
| The product worked                  | Post-trip: "I used this in a real conversation"             | > 60%     |
| Comprehension is the differentiator | Post-trip: hear-side drills rated more useful than say-side | > 50%     |

If the completion gate fails, the finish-line premise is wrong and the rest of the product does not matter.

## Appendix: Research Sources Referenced

- Duolingo Q1 and Q2 2026 results: 58.7M DAU, 137.8M MAU, 12.5M paid subscribers; 2026 strategy shift toward user growth and broader free access to AI features
- Duolingo Max feature set: Roleplay, Video Call, Adventures (cafe ordering, passport check scenarios)
- Babbel positioning and pricing: 14 languages, 10 to 15 minute lessons, travel-oriented marketing, roughly \$8 to \$15 per month
- Speak: \$162M raised, \$1B valuation, six languages, \$17.99/month
- Memrise: native-speaker video clips for real-speed listening, MemBot rehearsal
- Mondly: 41 languages, 50 practical topics across restaurant, hotel, transport, emergency
- TripLingo: activity-based custom phrase lists, 13 languages, 2,000+ phrases each, 30,000 audio files, 500,000+ travelers, \$2.3M raised, acquired 2019, consumer app dormant since 2021
- Education app benchmarks: 95%+ churn by day 30, roughly 2% day-30 retention
- Fodor's Travel Talk Forums: comprehension asymmetry, "too general for our needs"
- Duolingo content critique: opening-lesson vocabulary not ranked by travel frequency
- Regional vocabulary mismatch: abacaxi vs ananas in Portugal
- Phrasebook coverage gaps: Thai app missing "no," Arabic app missing the common greeting